

MACHINE LEARNING IN R: PACKAGE MLR


Bernd Bischl
Computational Statistics, LMU



AGENDA

- About mlr
- Features of mlr
 - ▶ Tasks and Learners
 - ▶ Train, Test, Resample
 - ▶ Benchmarking
 - ▶ Hyperparameter Tuning
 - ▶ Performance Visualization
- iml - Interpretable Machine Learning
- OpenML

MACHINE LEARNING



Machine Learning is a method of teaching computers to make predictions based on some data.

MOTIVATION

THE GOOD NEWS

- CRAN serves hundreds of packages for machine learning
- Often compliant to the unwritten interface definition:

```
> model = fit(target ~ ., data = train.data, ...)  
> predictions = predict(model, newdata = test.data, ...)
```

THE BAD NEWS

- Some packages API is “just different”
- Functionality is always package or model-dependent, even though the procedure might be general
- No meta-information available or buried in docs

Our goal: A domain-specific language for many machine learning concepts!

ABOUT

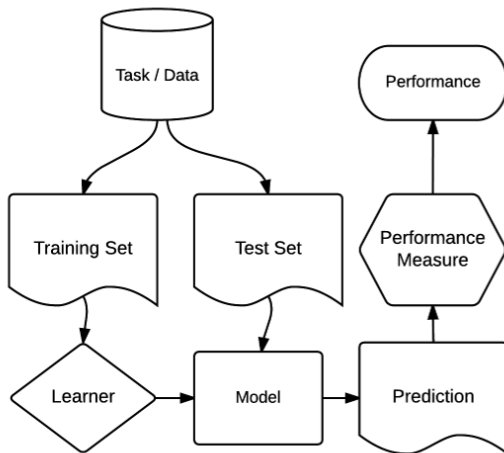
- Project home page

`https://github.com/mlr-org/mlr`

- ▶ [Cheatsheet](#) for an quick overview
 - ▶ [Tutorial](#) for mlr documentation with many code examples
 - ▶ Ask questions in the [GitHub issue tracker](#)
- 8-10 main developers, quite a few contributors, 4 GSOC projects in 2015/16 and one coming in 2017
 - About 20K lines of code, 8K lines of unit tests

MOTIVATION: MLR

- Unified interface for the basic building blocks: tasks, learners, hyperparameters, ...




R EXAMPLE

The mlr process

R EXAMPLE: TITANIC

Titanic: Machine Learning from Disaster

- Titanic sinking on April 15, 1912
- Data on Kaggle: <https://www.kaggle.com/c/titanic>
- 1502 out of 2224 passengers got killed
- Task
 - ▶ Can we predict who survived?
 - ▶ Why did people die / Which groups?



R EXAMPLE: DATA SET

■ Data Dictionary

Survived	Survived, 0 = No, 1 = Yes
Pclass	Ticket class, from 1st to 3rd
Sex	Sex
Age	Age in years
Sibsp	# of siblings/ spouses
Parch	# of parents/ children
Ticket	Ticket number
Fare	Passenger fare
Cabin	Cabin number
Embarked	Port of Embarkation

PREPROCESSING I

- Load the input data
- Combine training and test data

```
> train = read.table("train.csv", header = TRUE, sep = ",",  
+   colClasses = c("integer", "factor", "factor", "character",  
+   "factor", "numeric", "numeric", "numeric",  
+   "factor", "numeric", "factor", "factor"))  
> train$train = TRUE  
>  
> test = read.table("test.csv", header = TRUE, sep = ",",  
+   colClasses = c("integer", "factor", "character", "factor",  
+   "numeric", "numeric", "numeric", "factor", "numeric",  
+   "factor", "factor"))  
> test$Survived = NA  
> test$train = FALSE  
>  
> data = rbind(train, test)
```

PREPROCESSING II

- Set empty factor levels to NA

```
> data$Embarked[data$Embarked == ""] = NA  
> data$Embarked = droplevels(data$Embarked)  
> data$Cabin[data$Cabin == ""] = NA  
> data$Cabin = droplevels(data$Cabin)
```

PREPROCESSING III

```
> summarizeColumns(data)[, -c(5, 6,7)]
```

##		name	type	na	mean	min	max	nlevs
## 1	PassengerId	integer	0	655.0000	1.00	1309.0	0	
## 2	Survived	factor	418	NA	342.00	549.0	2	
## 3	Pclass	factor	0	NA	277.00	709.0	3	
## 4	Name	character	0	NA	1.00	2.0	1307	
## 5	Sex	factor	0	NA	466.00	843.0	2	
## 6	Age	numeric	263	29.8811	0.17	80.0	0	
## 7	SibSp	numeric	0	0.4989	0.00	8.0	0	
## 8	Parch	numeric	0	0.3850	0.00	9.0	0	
## 9	Ticket	factor	0	NA	1.00	11.0	929	
## 10	Fare	numeric	1	33.2955	0.00	512.3	0	
## 11	Cabin	factor	1014	NA	1.00	6.0	186	
## 12	Embarked	factor	2	NA	123.00	914.0	3	
## 13	train	logical	0	NA	418.00	891.0	2	

PREPROCESSING I

```
> # Price per person, multiple tickets bought by one person
> data$farePp = data$Fare / (data$Parch + data$SibSp + 1)
>
> # The deck can be extracted from the the cabin number
> data$deck = as.factor(str_sub(data$Cabin, 1, 1))
>
> # Starboard had an odd number, portside even cabin numbers
> data$portside = str_sub(data$Cabin, 3, 3)
> data$portside = as.numeric(data$portside) %% 2

## Warning:  NAs introduced by coercion

> # Drop stuff we cannot easily model on
> data = dropNamed(data,
+   c("Cabin", "PassengerId", "Ticket", "Name"))
```

IMPUTATION

- Remove missing values
- Impute numerics with median and factors with a separate category
- NB: This is really naive and we should probably use multiple imputation

```
> data = impute(data, cols = list(  
+   Age = imputeMedian(),  
+   Fare = imputeMedian(),  
+   Embarked = imputeConstant("__miss__"),  
+   farePp = imputeMedian(),  
+   deck = imputeConstant("__miss__"),  
+   portside = imputeConstant("__miss__")  
+ ))  
>  
> data = data$data  
> data = convertDataFrameCols(data, chars.as.factor = TRUE)
```

TASKS I

- Split back into train and test
- Create classification problem

```
> dtrain = data[data$train, ]  
> dtrain$train = NULL  
> dtest = data[!data$train, ]  
> dtest$train = NULL  
>  
> task = makeClassifTask(id = "titanic", data = dtrain,  
+   target = "Survived", positive = "1")
```

TASKS II

```
> print(task)

## Supervised task: titanic
## Type: classif
## Target: Survived
## Observations: 891
## Features:
##      numerics      factors      ordered functionals
##           5           5           0           0
## Missings: FALSE
## Has weights: FALSE
## Has blocking: FALSE
## Has coordinates: FALSE
## Classes: 2
##    0    1
## 549 342
## Positive class: 1
```


WHAT LEARNERS ARE AVAILABLE? I

CLASSIFICATION (84)

- LDA, QDA, RDA, MDA
- Trees and forests
- Boosting (different variants)
- SVMs (different variants)
- ...

CLUSTERING (9)

- K-Means
- EM
- DBscan
- X-Means
- ...

REGRESSION (61)

- Linear, lasso and ridge
- Boosting
- Trees and forests
- Gaussian processes
- ...

SURVIVAL (12)

- Cox-PH
- Cox-Boost
- Random survival forest
- Penalized regression
- ...

WHAT LEARNERS ARE AVAILABLE? II

- Explore all learners via tutorial
- Or ask mlr

```
> listLearners("classif",  
+   properties = c("prob", "multiclass"))[1:5, c(1,4,13,16)]
```

##	class	package	prob	multiclass
## 1	classif.adaboostm1	RWeka	TRUE	TRUE
## 2	classif.boosting	adabag,rpart	TRUE	TRUE
## 3	classif.C50	C50	TRUE	TRUE
## 4	classif.cforest	party	TRUE	TRUE
## 5	classif.ctree	party	TRUE	TRUE

TRAIN - TEST - EVAL I

```
> lrn = makeLearner("classif.rpart", predict.type = "prob")
> mod = train(lrn, task)
> pred = predict(mod, newdata = dtest)
> head(as.data.frame(pred))
```

##	truth	prob.0	prob.1	response
## 892	<NA>	0.8981	0.1019	0
## 893	<NA>	0.8571	0.1429	0
## 894	<NA>	0.8981	0.1019	0
## 895	<NA>	0.8981	0.1019	0
## 896	<NA>	0.3297	0.6703	1
## 897	<NA>	0.8981	0.1019	0

TRAIN - TEST - EVAL II

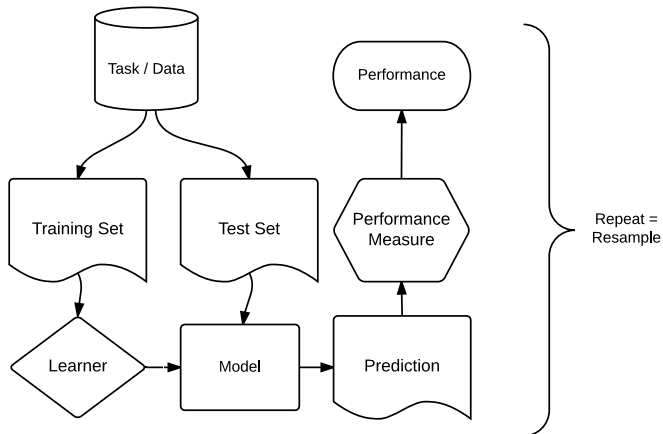
- We don't have labels for the true test set here
- Let's eval on train set now
- Which is dangerous in practice and not recommended in general!

```
> pred = predict(mod, newdata = dtrain)
> performance(pred, measures = list(mlr::acc, mlr::auc))

##      acc      auc
## 0.8541 0.8843
```

RESAMPLING

- Aim: Assess the performance of a learning algorithm
- Uses the data more efficiently than simple train-test
- Repeatedly split in train and test, then aggregate results.



CROSS VALIDATION I

- Most popular resampling strategy: Cross validation with 5 or 10 folds
- Split the data into k roughly equally-sized partitions
- Use each part once as test set and joint $k - 1$ other parts to train
- Obtain k test errors and average them

Example of 3-fold cross-validation

Iteration 1	Test	Train	Train
Iteration 2	Train	Test	Train
Iteration 3	Train	Train	Test

CROSSVALIDATION IN MLR I

```
> lrn = makeLearner("classif.rpart", predict.type = "prob")
> rdesc = makeResampleDesc("CV", iters = 3, stratify = TRUE)
> r = resample(lrn, task, rdesc,
+   measures = list(mlr::acc, mlr::auc))
> print(r)

## Resample Result
## Task: titanic
## Learner: classif.rpart
## Aggr perf: acc.test.mean=0.8013,auc.test.mean=0.8439
## Runtime: 0.0378931
```

CROSSVALIDATION IN MLR II

```
> head(r$measures.test)
```

```
##      iter      acc      auc  
## 1      1 0.8148 0.8184  
## 2      2 0.8182 0.8667  
## 3      3 0.7710 0.8467
```

```
> head(as.data.frame(r$pred))
```

```
##      id truth prob.0 prob.1 response iter  set  
## 1   8      0 0.8831 0.1169         0    1 test  
## 2  15      0 0.1613 0.8387         1    1 test  
## 3  17      0 0.8831 0.1169         0    1 test  
## 4  21      0 0.8831 0.1169         0    1 test  
## 5  25      0 0.2500 0.7500         1    1 test  
## 6  27      0 0.8831 0.1169         0    1 test
```


RESAMPLING IN MLR

- Holdout (Train-Test): “Holdout”
- Cross Validation: “CV”
- Leave-one-out: “LOO”
- Subsample (Monte-Carlo CV) “Subsample”
- Out-of-bag bootstrap and other methods “Bootstrap”

BENCHMARKING AND MODEL COMPARISON I

- Comparison of multiple models on multiple data sets
- Aim: Find best learners for a data set or domain, learn about learner characteristics, ...

```
> bmr = benchmark(list.of.learners, list.of.tasks, rdesc)
```


R EXAMPLE: ALGORITHMS I

■ Benchmark experiment - Compare 4 algorithms

```
> set.seed(3)
> learners = c("glmnet", "naiveBayes", "randomForest", "ksvm")
> learners = makeLearners(learners, type = "classif",
+   predict.type = "prob")
> bmr = benchmark(learners, task, rinst, measures = mlr::auc)
```

R EXAMPLE: ALGORITHMS II

```
> plotBMRBoxplots(bmr)
```




HYPERPARAMETER TUNING

- Aim: Optimize parameters or decisions for an machine learning algorithm w.r.t. the estimated prediction error
- Used to find “best” hyperparameters for a method in a data-dependent way
- General procedure: Tuner proposes param point, eval by resampling, feedback value to tuner

GRID SEARCH

Try all combinations of finite grid


↪ Inefficient, combinatorial explosion, searches irrelevant areas



RANDOM SEARCH

Uniformly randomly draw configurations,

↪ Scales better than grid search, easily extensible



TUNING IN MLR I

- Create a set of parameters

```
> lrn = makeLearner("classif.ksvm", predict.type = "prob")
> par.set = makeParamSet(
+   makeNumericParam("C", lower = -8, upper = 8,
+     trafo = function(x) 2^x),
+   makeNumericParam("sigma", lower = -8, upper = 8,
+     trafo = function(x) 2^x)
+ )
```


TUNING IN MLR II

- Optimize the hyperparameter of learner

```
> ctrl = makeTuneControlGrid(resolution = 7)
> tr = tuneParams(lrn, task = task, par.set = par.set,
+   resampling = rdesc, control = ctrl,
+   measures = mlr::auc)
```

TUNING IN MLR III

```
> head(as.data.frame(tr$opt.path))[, c(1,2,3,7)]
```

##		C	sigma	auc.test.mean	exec.time
## 1	-8.000	-8		0.8143	0.385
## 2	-5.333	-8		0.8152	0.397
## 3	-2.667	-8		0.8150	0.408
## 4	0.000	-8		0.8375	0.381
## 5	2.667	-8		0.8449	0.359
## 6	5.333	-8		0.8413	0.368


R EXAMPLE: TUNING I

- We used all algorithms in their default settings
- Hopefully tuning will improve the performance
- Nested cross validation to get true out-of-sample predictions

```
> par.set = makeParamSet(  
+   makeNumericParam("C", lower = -8, upper = 8,  
+     trafo = function(x) 2^x),  
+   makeNumericParam("sigma", lower = -8, upper = 8,  
+     trafo = function(x) 2^x)  
+ )  
> tune.ctrl = makeTuneControlRandom(maxit = 10L)  
> classif.ksvm.tuned = makeTuneWrapper(learners$classif.ksvm,  
+   resampling = cv3, par.set = par.set, control = tune.ctrl)  
> bmr2 = benchmark(classif.ksvm.tuned, task.train, rinst)
```

R EXAMPLE: TUNING II

```
> plotBMRBoxplots(mergeBenchmarkResults(list(bmr, bmr2)))
```



PERFORMANCE MEASURES

- Different performance measures for different types of learning problems
- In `mlr` you can check out all implemented measures via `https://mlr-org.github.io/mlr/articles/tutorial/devel/measures.html`





PERFORMANCE MEASURE FOR CLASSIFICATION I

- In our Titanic example we have a classification problem
- Confusion matrix:
contingency table of predictions \hat{y} and true labels y

Diagnostic Testing Measures				
		Actual Class y		
		Positive	Negative	
\hat{y} Test outcome	Test outcome positive	True positive (TP)	False positive (FP, Type I error)	Precision = $\frac{\#TP}{\#TP + \#FP}$
	Test outcome negative	False negative (FN, Type II error)	True negative (TN)	Negative predictive value = $\frac{\#TN}{\#FN + \#TN}$
		Sensitivity = $\frac{\#TP}{\#TP + \#FN}$	Specificity = $\frac{\#TN}{\#FP + \#TN}$	Accuracy = $\frac{\#TP + \#TN}{\#TOTAL}$

PERFORMANCE MEASURE FOR CLASSIFICATION II

- For classification performance measure the True Positive Rate (TPR) and the False Positive Rate (FPR) are plotted → ROC Curve (Receiver Operating Characteristic)




- For measuring the performance we can calculate the area under the ROC curve (AUC)

R EXAMPLE: RANDOM FOREST I

- The Random Forest seems to work best, lets have a closer look

```
> r = holdout(learners$classif.randomForest, task)
> df = generateThreshVsPerfData(r$pred, list(fpr, tpr, acc))
> plotROCCurves(df)
```


R EXAMPLE: RANDOM FOREST II



R EXAMPLE: RANDOM FOREST III

```
> calculateROCMeasures(pred)

##      predicted
## true 0      1
##    0 504      45      tpr: 0.75 fnr: 0.25
##    1 85      257      fpr: 0.08 tnr: 0.92
##      ppv: 0.85 for: 0.14 lrp: 9.17 acc: 0.85
##      fdr: 0.15 npv: 0.86 lrm: 0.27 dor: 33.86
##
##
## Abbreviations:
## tpr - True positive rate (Sensitivity, Recall)
## fpr - False positive rate (Fall-out)
## fnr - False negative rate (Miss rate)
## tnr - True negative rate (Specificity)
## ppv - Positive predictive value (Precision)
## for - False omission rate
## lrp - Positive likelihood ratio (LR+)
## fdr - False discovery rate
## npv - Negative predictive value
## acc - Accuracy
```

PARALLELIZATION

- We use our own package: `parallelMap`
- Setup:

```
> parallelStart("multicore")  
> benchmark(...)  
> parallelStop()
```

- Backends: `local`, `multicore`, `socket`, `mpi` and `batchtools`
- The latter means support for: makeshift SSH-clusters, Docker swarm and HPC schedulers like SLURM, Torque/PBS, SGE or LSF
- Levels allow fine grained control over the parallelization
 - ▶ `mlr.resample`: Job = "train / test step"
 - ▶ `mlr.tuneParams`: Job = "resample with these parameter settings"
 - ▶ `mlr.selectFeatures`: Job = "resample with this feature subset"
 - ▶ `mlr.benchmark`: Job = "evaluate this learner on this data set"

INTERPRETABLE MACHINE LEARNING

- `iml` - Interpretable Machine Learning -
<https://github.com/christophM/iml>
- Background
 - ▶ Machine learning has a huge potential
 - ▶ Lack of explanation hurts trusts and creates barrier for machine learning adoption
 - ▶ Interpretation of the behaviour and explanation of predictions of machine learning model with **Interpretable Machine Learning**

SUPPORTED METHODS

- Model-agnostic interpretability methods for **any** kind of machine learning model
- Supported are
 - ▶ Feature importance
 - ▶ Partial dependence plots
 - ▶ Individual conditional expectation plots
 - ▶ Tree surrogate
 - ▶ Local interpretable model-agnostic explanations
 - ▶ Shapley value

ONE IML MODEL FOR ALL METHODS

- Use iml package

```
> library(impl)
```

- We use our trained model mod

```
> mod
```

```
## Model for learner.id=classif.randomForest; learner.class=clas  
## Trained on: task.id = titanic; obs = 594; features = 12  
## Hyperparameters:
```


- Extract features
- Create IML model

```
> X = dropNamed(dtrain, "Survived")  
> impl.mod = Predictor$new(mod, data = X,  
+   y = train$Survived, class = 2)
```

FEATURE IMPORTANCE

- What were the most important features?

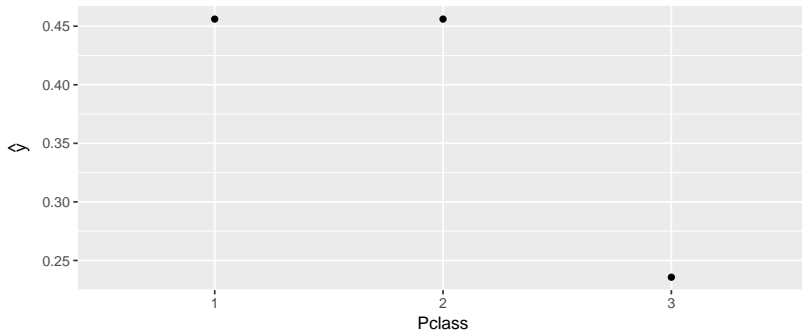
```
> imp = FeatureImp$new(iml.mod, loss = "ce")  
> plot(imp)
```



PARTIAL DEPENDENCE PLOTS

- How does the “passenger class” influence the prediction on average?


```
> pdp = PartialDependence$new(iml.mod, feature = "Pclass")  
> plot(pdp)
```



LOCAL LINEAR MODELS (LIME)

- Explain a single prediction with LIME

```
> X[1,]  
  
##    Pclass  Sex Age SibSp Parch Fare Embarked farePp    deck p  
## 1      3 male  22     1     0  7.25          S   3.625 __miss__ _  
  
> lime = LocalModel$new(iml.mod, x.interest = X[1,])  
> plot(lime)
```



THERE IS MORE ...

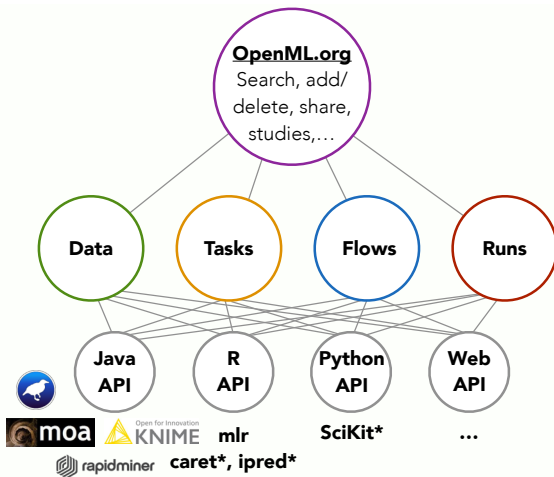
- Regression, Clustering and Survival analysis
- Cost-sensitive learning
- Multi-Label learning
- Imbalancy correction
- Wrappers
- Bayesian optimization
- Multi-criteria optimization
- Ensembles, generic bagging and stacking
- ...

WE ARE WORKING ON

- Even better tuning system
- More interactive and 3D plots
- Large-Scale learning on databases
- Time-Series tasks
- Large-Scale usage of OpenML
- `auto-mlr`
- ...

OPENML

Main idea: Make ML experiments reproducible, computer-readable and allow collaboration with others.



OPENML R-PACKAGE

<https://github.com/openml/r>

TUTORIAL

- Caution: Work in progress

CURRENT API IN R

- Explore and Download data and tasks
- Register learners and upload runs
- Explore your own and other people's results

MLR CONTRIBUTION

- Write an issue on Git
- We are founding an association - **Machine Learning in R e.V**
subscribe for updates contact.mlr.org@gmail.com

Thanks!