

# MACHINE LEARNING IN R: PACKAGE MLR

Bernd Bischl  
Computational Statistics, LMU




# WORKSHOP DOCUMENTATION

`goo.gl/DYzSmA`

# AGENDA

- About mlr
- Features of mlr
  - ▶ Tasks and Learners
  - ▶ Train, Test, Resample
  - ▶ Benchmarking
  - ▶ Hyperparameter Tuning
  - ▶ Nested Resampling
  - ▶ Performance Visualization
  - ▶ Parallelization
- iml - Interpretable Machine Learning
- mlrMBO - Bayesian Optimization
- mlrCPO - Composable Preprocessing
- OpenML
- Outlook and mlr contribution

# MACHINE LEARNING



Machine Learning is a method of teaching computers to make predictions based on some data.

# MOTIVATION

## THE GOOD NEWS

- CRAN serves hundreds of packages for machine learning
- Often compliant to the unwritten interface definition:

```
> model = fit(target ~ ., data = train.data, ...)  
> predictions = predict(model, newdata = test.data, ...)
```

## THE BAD NEWS

- Some packages API is “just different”
- Functionality is always package or model-dependent, even though the procedure might be general
- No meta-information available or buried in docs

Our goal: A domain-specific language for many machine learning concepts!

# ABOUT

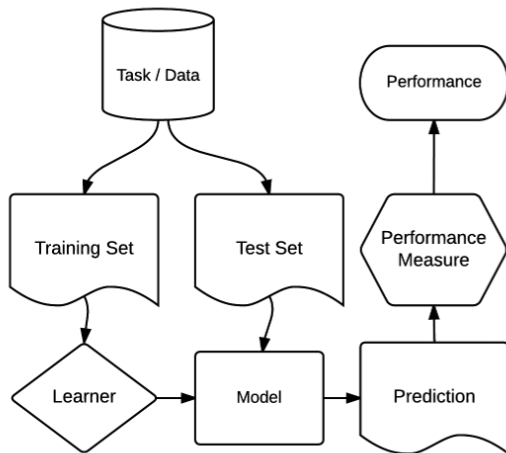
- Project home page

`https://github.com/mlr-org/mlr`

- ▶ Cheatsheet for an quick overview
  - ▶ Tutorial for mlr documentation with many code examples
  - ▶ Ask questions in the GitHub issue tracker
- 8-10 main developers, quite a few contributors, 4 GSOC projects in 2015/16 and one coming in 2017
  - About 20K lines of code, 8K lines of unit tests

# MOTIVATION: MLR

- Unified interface for the basic building blocks: tasks, learners, hyperparameters, ...



# R EXAMPLE


The mlr process



# R EXAMPLE: TITANIC

## Titanic: Machine Learning from Disaster

- Titanic sinking on April 15, 1912
- Data provided on our website [goo.gl/DYzSmA](http://goo.gl/DYzSmA)
- 809 out of 1309 passengers got killed
- Task
  - ▶ Can we predict who survived?
  - ▶ Why did people die / Which groups?



# R EXAMPLE: DATA SET

## ■ Data Dictionary

Survived	Survived, 0 = No, 1 = Yes
Pclass	Ticket class, from 1st to 3rd
Sex	Sex
Age	Age in years
Sibsp	# of siblings/ spouses
Parch	# of parents/ children
Ticket	Ticket number
Fare	Passenger fare
Cabin	Cabin number
Embarked	Port of Embarkation

# PREPROCESSING I

## ■ Load the input data

```
> load("data.rda")
> print(summarizeColumns(data)[, -c(5, 6, 7)], digits = 0)
```

##	name	type	na	mean	min	max	nlevs
## 1	Pclass	factor	0	NA	277	709	3
## 2	Survived	factor	0	NA	500	809	2
## 3	Name	character	0	NA	1	2	1307
## 4	Sex	factor	0	NA	466	843	2
## 5	Age	numeric	263	30	0	80	0
## 6	Sibsp	numeric	0	0	0	8	0
## 7	Parch	numeric	0	0	0	9	0
## 8	Ticket	factor	0	NA	1	11	929
## 9	Fare	numeric	1	33	0	512	0
## 10	Cabin	factor	0	NA	1	1014	187
## 11	Embarked	factor	0	NA	2	914	4

# PREPROCESSING II

- NB: All preprocessing steps are really naive, later we show better preprocessing with `mlrCP0`
- Set empty factor levels to NA

```
> data$Embarked[data$Embarked == ""] = NA  
> data$Embarked = droplevels(data$Embarked)  
> data$Cabin[data$Cabin == ""] = NA  
> data$Cabin = droplevels(data$Cabin)
```

# PREPROCESSING III

```
> # Price per person, multiple tickets bought by one
> # person
> data$farePp = data$Fare / (data$Parch + data$Sibsp + 1)
>
> # The deck can be extracted from the the cabin number
> data$deck = as.factor(stri_sub(data$Cabin, 1, 1))
>
> # Starboard had an odd number, portside even cabin
> # numbers
> data$portside = stri_sub(data$Cabin, 3, 3)
> data$portside = as.numeric(data$portside) %% 2
>
> # Drop stuff we cannot easily model on
> data = dropNamed(data,
+   c("Cabin", "PassengerId", "Ticket", "Name"))
```

# PREPROCESSED DATA

```
> print(summarizeColumns(data)[, -c(5, 6, 7)], digits = 0)
```

##	name	type	na	mean	min	max	nlevs
## 1	Pclass	factor	0	NA	277	709	3
## 2	Survived	factor	0	NA	500	809	2
## 3	Sex	factor	0	NA	466	843	2
## 4	Age	numeric	263	30	0	80	0
## 5	Sibsp	numeric	0	0	0	8	0
## 6	Parch	numeric	0	0	0	9	0
## 7	Fare	numeric	1	33	0	512	0
## 8	Embarked	factor	2	NA	123	914	3
## 9	farePp	numeric	1	21	0	512	0
## 10	deck	factor	1014	NA	1	94	8
## 11	portside	numeric	1059	0	0	1	0

# IMPUTATION

- Remove missing values
- Impute numerics with median and factors with a separate category
- NB: This is really naive, we should probably use multiple imputation and embed this in cross-validation

```
> data = impute(data, cols = list(  
+   Age = imputeMedian(),  
+   Fare = imputeMedian(),  
+   Embarked = imputeConstant("__miss__"),  
+   farePp = imputeMedian(),  
+   deck = imputeConstant("__miss__"),  
+   portside = imputeConstant("__miss__")  
+ ))  
>  
> data = data$data  
> data = convertDataFrameCols(data, chars.as.factor = TRUE)
```

# TASKS I

- Create classification problem

```
> task = makeClassifTask(id = "titanic", data = data,  
+   target = "Survived", positive = "1")
```



# TASKS II

```
> print(task)

## Supervised task: titanic
## Type: classif
## Target: Survived
## Observations: 1309
## Features:
##      numerics      factors      ordered functionals
##           5           5           0           0
## Missings: FALSE
## Has weights: FALSE
## Has blocking: FALSE
## Has coordinates: FALSE
## Classes: 2
##    0    1
## 809 500
## Positive class: 1
```

# WHAT LEARNERS ARE AVAILABLE? I

## CLASSIFICATION (84)

- LDA, QDA, RDA, MDA
- Trees and forests
- Boosting (different variants)
- SVMs (different variants)
- ...

## CLUSTERING (9)

- K-Means
- EM
- DBscan
- X-Means
- ...

## REGRESSION (61)


- Linear, lasso and ridge
- Boosting
- Trees and forests
- Gaussian processes
- ...

## SURVIVAL (12)

- Cox-PH
- Cox-Boost
- Random survival forest
- Penalized regression
- ...

# WHAT LEARNERS ARE AVAILABLE? II

- Explore all learners via tutorial

 2.13 Get Started Basics ▾ Advanced ▾ Extending ▾ Appendix ▾ mlr-org Packages ▾ <input data-bbox="1136 253 1260 315" type="text" value="Search..."/>									
Class / Short Name / Name	Packages	Num.	Fac.	Ord.	NAs	Weights	Props	Note	
<b>classif.ada</b> <i>ada</i>  ada Boosting	<a href="#">ada</a> <a href="#">rpart</a>	X	X				prob twoclass	xval has been set to 0 by default for spec	
<b>classif.adaboostm1</b> <i>adaboostm1</i>  ada Boosting M1	<a href="#">RWeka</a>	X	X				prob twoclass multiclass	NAs are directly passed to WEKA with na.ac	
<b>classif.bartMachine</b> <i>bartmachine</i>  Bayesian Additive Regression Trees	<a href="#">bartMachine</a>	X	X		X		prob twoclass	use_missing_data has been set to TRUE	
<b>classif.binomial</b> <i>binomial</i>  Binomial Regression	<a href="#">stats</a>	X	X			X	prob twoclass	Delegates to glm with freely choosable bini	

# WHAT LEARNERS ARE AVAILABLE? III

## ■ Or ask mlr

```
> listLearners("classif", properties = c("prob",  
+   "multiclass"))[1:5, c(1,4,13,16)]
```

##	class	package	prob	multiclass
## 1	classif.adaboostm1	RWeka	TRUE	TRUE
## 2	classif.boosting	adabag,rpart	TRUE	TRUE
## 3	classif.C50	C50	TRUE	TRUE
## 4	classif.cforest	party	TRUE	TRUE
## 5	classif.ctree	party	TRUE	TRUE

# TRAIN MODEL I

- Create a learner
- Output posterior probs – instead of a factor of class labels

```
> lrn = makeLearner("classif.randomForest",  
+   predict.type = "prob")
```

- Split data into a training and test data set (necessary for performance evaluation)
- And train a model

```
> n = nrow(data)  
> train = sample(n, size = 2/3 * n)  
> test = setdiff(1:n, train)  
>  
> mod = train(lrn, task, subset = train)
```

# PREDICTIONS I

- Make predictions for new data

```
> pred = predict(mod, task = task, subset = test)
> head(as.data.frame(pred))
```

##	id	truth	prob.0	prob.1	response
## 3	3	0	0.500	0.500	1
## 8	8	0	0.814	0.186	0
## 9	9	1	0.016	0.984	1
## 12	12	1	0.176	0.824	1
## 13	13	1	0.018	0.982	1
## 15	15	1	0.868	0.132	0

# PREDICTIONS II

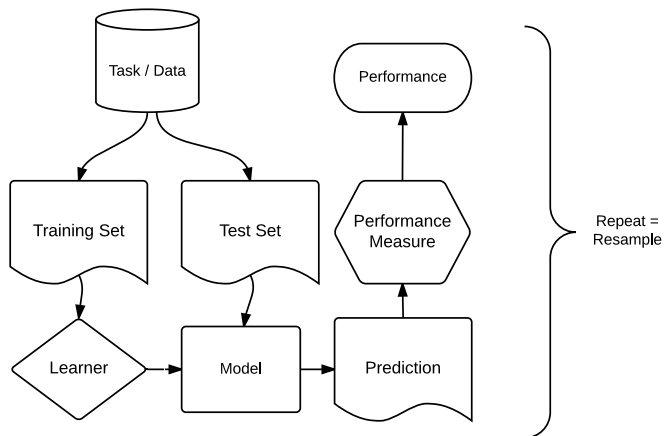
- Evaluate predictive performance

```
> performance(pred, measures = list(mlr::acc, mlr::auc))
```

```
##          acc          auc  
## 0.8169336 0.8760819
```

# RESAMPLING

- Aim: Assess the performance of a learning algorithm
- Uses the data more efficiently than simple train-test
- Repeatedly split in train and test, then aggregate results.





# CROSS VALIDATION

- Most popular resampling strategy: Cross validation with 5 or 10 folds
- Split the data into  $k$  roughly equally-sized partitions
- Use each part once as test set and joint  $k - 1$  other parts to train
- Obtain  $k$  test errors and average them

## Example of 3-fold cross-validation

Iteration 1	Test	Train	Train
Iteration 2	Train	Test	Train
Iteration 3	Train	Train	Test

# CROSSVALIDATION IN MLR I

```
> rdsc = makeResampleDesc("CV", iters = 3,  
+   stratify = TRUE)  
>  
> r = resample(lrn, task, rdsc,  
+   measures = list(mlr::acc, mlr::auc))  
> print(r)  
  
## Resample Result  
## Task: titanic  
## Learner: classif.randomForest  
## Aggr perf: acc.test.mean=0.7906791,auc.test.mean=0.8548651  
## Runtime: 4.30436
```

# CROSSVALIDATION IN MLR II

```
> head(r$measures.test)
```

##	iter	acc	auc
## 1	1	0.7917620	0.8678310
## 2	2	0.7912844	0.8518965
## 3	3	0.7889908	0.8448679

```
> head(as.data.frame(r$pred))
```

##	id	truth	prob.0	prob.1	response	iter	set
## 1	3	0	0.522	0.478	0	1	test
## 2	4	0	0.462	0.538	1	1	test
## 3	10	0	0.926	0.074	0	1	test
## 4	26	0	0.566	0.434	0	1	test
## 5	35	0	0.448	0.552	1	1	test
## 6	39	0	0.402	0.598	1	1	test

# RESAMPLING METHODS IN MLR

Methods	Parameter
<b>CV</b>	iters stratify
<b>LOO</b>	
<b>RepCV</b>	reps folds stratify
<b>Bootstrap</b>	iters stratify
<b>Subsample</b>	iters split stratify
<b>Holdout</b>	split stratify

# BENCHMARKING AND MODEL COMPARISON I

- Comparison of multiple models on multiple data sets
- Aim: Find best learners for a data set or domain, learn about learner characteristics, ...

```
> bmr = benchmark(list.of.learners, list.of.tasks, rdesc)
```

# R EXAMPLE: ALGORITHMS I

## ■ Benchmark experiment - Compare 4 algorithms

```
> set.seed(3)
>
> learners = c("glmnet", "naiveBayes", "randomForest",
+   "ksvm")
> learners = makeLearners(learners, type = "classif",
+   predict.type = "prob")
>
> bmr = benchmark(learners, task, rdesc,
+   measures = mlr::auc)
```

# R EXAMPLE: ALGORITHMS II

## ■ Access aggregated results

```
> getBMRAggrPerformances(bmr, as.df = TRUE)
```

##	task.id	learner.id	auc.test.mean
## 1	titanic	classif.glmnet	0.8402273
## 2	titanic	classif.naiveBayes	0.8011408
## 3	titanic	classif.randomForest	0.8571534
## 4	titanic	classif.ksvm	0.8292053

## R EXAMPLE: ALGORITHMS III

- Access more fine-grained results
- Many more getters for predictions, models, etc.


```
> head(getBMRPerformances(bmr, as.df = TRUE), 4)
```

##	task.id	learner.id	iter	auc
## 1	titanic	classif.glmnet	1	0.8378909
## 2	titanic	classif.glmnet	2	0.8136701
## 3	titanic	classif.glmnet	3	0.8691209
## 4	titanic	classif.naiveBayes	1	0.8006653



# R EXAMPLE: ALGORITHMS IV

```
> plotBMRBoxplots(bmr)
```



# PERFORMANCE MEASURES

- Different performance measures for different types of learning problems
- mlr has 71 performance measures implemented
- See all via <https://mlr-org.github.io/mlr/articles/tutorial/devel/measures.html> or `listMeasures()`

mlr 2.13 Get Started Basics ▾ Advanced ▾ Extending ▾ Appendix ▾ mlr-org Packages ▾ Search...												
Classification												
ID / Name	Minim.	Best	Worst	Multi	Pred.	Truth	Probs	Model	Task	Feats	Aggr.	Note
<a href="#">acc</a> Accuracy		1	0	X	X	X					<a href="#">test_mean</a>	Defined as: $\text{mean}(\text{response} == \text{truth})$
<a href="#">auc</a> Area under the curve		1	0		X	X	X				<a href="#">test_mean</a>	Integral over the graph that results fr and tpr for many different thresholds
<a href="#">bac</a> Balanced accuracy		1	0	X	X	X					<a href="#">test_mean</a>	For binary tasks, mean of true positive negative rate.
<a href="#">ber</a> Balanced error rate	X	0	1	X	X	X					<a href="#">test_mean</a>	Mean of misclassification error rates classes.
<a href="#">brier</a> Brier score	X	0	1		X	X	X				<a href="#">test_mean</a>	The Brier score is defined as the qua between the probability and the valu That means we use the numeric repr for our target classes. It is similar to error in regression. multiclass.brier is one vs. all comparisons and for a bin brier.
<a href="#">brier.scaled</a> Brier scaled		1	0		X	X	X				<a href="#">test_mean</a>	Brier score scaled to [0,1], see <a href="http://www.ncbi.nlm.nih.gov/omc/ar">http://www.ncbi.nlm.nih.gov/omc/ar</a>
<a href="#">f1</a> F1 measure		1	0		X	X					<a href="#">test_mean</a>	Defined as: $2 * \text{tp} / (\text{sum}(\text{truth} == \text{posi sum}(\text{response} == \text{positive}))$





# PERFORMANCE MEASURE FOR CLASSIFICATION I

- In our Titanic example we have a classification problem
- Confusion matrix:  
contingency table of predictions  $\hat{y}$  and true labels  $y$

Diagnostic Testing Measures				
		Actual Class $y$		
		Positive	Negative	
$\hat{y}$ Test outcome	Test outcome positive	True positive (TP)	False positive (FP, Type I error)	Precision = $\frac{\#TP}{\#TP + \#FP}$
	Test outcome negative	False negative (FN, Type II error)	True negative (TN)	Negative predictive value = $\frac{\#TN}{\#FN + \#TN}$
		Sensitivity = $\frac{\#TP}{\#TP + \#FN}$	Specificity = $\frac{\#TN}{\#FP + \#TN}$	Accuracy = $\frac{\#TP + \#TN}{\#TOTAL}$

# PERFORMANCE MEASURE FOR CLASSIFICATION II

- For classification performance measure the True Positive Rate (TPR) and the False Positive Rate (FPR) are plotted → ROC Curve (Receiver Operating Characteristic)




- For measuring the performance we can calculate the area under the ROC curve (AUC)

# R EXAMPLE: RANDOM FOREST I

- The Random Forest seems to work best, lets have a closer look

```
> res = holdout(lrn, task)
> df = generateThreshVsPerfData(res$pred,
+   list(fpr, tpr, acc))
> plotROCCurves(df)
```

## R EXAMPLE: RANDOM FOREST II




## R EXAMPLE: RANDOM FOREST III

```
> print(calculateROCMeasures(pred), abbreviations = FALSE)

##      predicted
## true 0        1
##    0 238      30      tpr: 0.7  fnr: 0.3
##    1 50       119     fpr: 0.11 tnr: 0.89
##      ppv: 0.8 for: 0.17 lrp: 6.29 acc: 0.82
##      fdr: 0.2 npv: 0.83 lrm: 0.33 dor: 18.88
```

# HYPERPARAMETER TUNING

- Optimize parameters or decisions for ML algorithm w.r.t. the estimated prediction error
- Tuner proposes configuration, eval by resampling, tuner receives performance, iterate






# GRID SEARCH

Try all combinations of finite grid


↪ Inefficient, combinatorial explosion, searches irrelevant areas



# RANDOM SEARCH

Uniformly randomly draw configurations,

↪ Scales better than grid search, easily extensible



# TUNING IN MLR I

- Create a set of parameters
- Here we optimize an RBF SVM on logscale

```
> lrn.ksvm = makeLearner("classif.ksvm",  
+   predict.type = "prob")  
>  
> par.set = makeParamSet(  
+   makeNumericParam("C", lower = -8, upper = 8,  
+     trafo = function(x) 2^x),  
+   makeNumericParam("sigma", lower = -8, upper = 8,  
+     trafo = function(x) 2^x)  
+ )
```

# TUNING IN MLR II

- Optimize the hyperparameter of learner

```
> tune.ctrl = makeTuneControlRandom(maxit = 10L)
> tr = tuneParams(lrn.ksvm, task = task, par.set = par.set,
+   resampling = rdesc, control = tune.ctrl,
+   measures = mlr::auc)
```

# TUNING IN MLR III

```
> head(as.data.frame(tr$opt.path))[, c(1,2,3,7)]
```

##		C	sigma	auc.test.mean	exec.time
## 1	7.803771	2.0060031	0.7570655	2.54	
## 2	-4.374242	-0.3324129	0.8160881	0.82	
## 3	-5.417617	3.5509443	0.7770489	0.89	
## 4	2.076026	-1.9390989	0.8135859	0.77	
## 5	1.887830	-4.4571549	0.8321945	0.82	
## 6	-2.167479	2.1372494	0.7861856	0.86	

# R EXAMPLE: TUNING I


- We used all algorithms in their default settings
- Hopefully tuning will improve the performance
- Nested cross validation to get true out-of-sample predictions

```
> classif.ksvm.tuned = makeTuneWrapper(  
+   lrn.ksvm, resampling = rdesc,  
+   par.set = par.set, control = tune.ctrl)  
> bmr2 = benchmark(classif.ksvm.tuned, task, rdesc)
```

- `makeTuneWrapper`: Fuses a base learner with a search strategy to select its hyperparameters

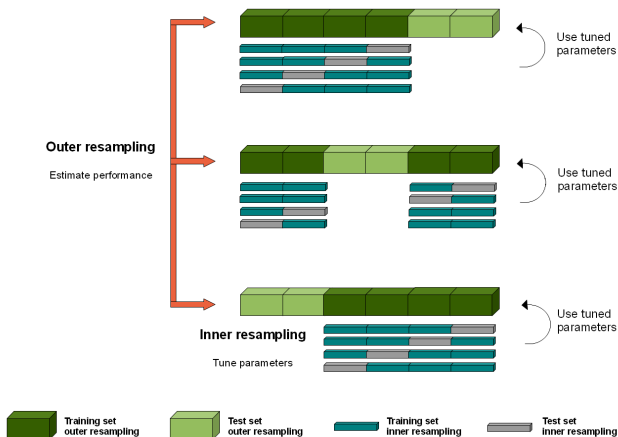
## R EXAMPLE: TUNING II

```
> plotBMRBoxplots(mergeBenchmarkResults(list(bmr, bmr2)))
```



# NESTED RESAMPLING I

- Danger of overfitting if hyperparameter and performance evaluation on the same test set
- Solution: 3-way split into training, optimization and test set





# NESTED RESAMPLING EXAMPLE I

- We show nested resampling with tuning
- Therefore we need an additional inner resampling loop

```
> inner.rdesc = makeResampleDesc("Subsample", iters = 2)
>
> classif.ksvm.inner = makeTuneWrapper(
+   lrn.ksvm, resampling = inner.rdesc,
+   par.set = par.set, control = tune.ctrl,
+   measures = mlr::auc)
```

# NESTED RESAMPLING EXAMPLE II

- We use `rdesc` for the outer loop

```
> r.nest = resample(classif.ksvm.inner, task,  
+   resampling = rdesc, extract = getTuneResult,  
+   measures = mlr::auc)  
> r.nest  
  
## Resample Result  
## Task: titanic  
## Learner: classif.ksvm.tuned  
## Aggr perf: auc.test.mean=0.8373767  
## Runtime: 13.2617
```

# NESTED RESAMPLING EXAMPLE III

```
> r.nest$extract

## [[1]]
## Tune result:
## Op. pars: C=3.39; sigma=0.0317
## auc.test.mean=0.8448094
##
## [[2]]
## Tune result:
## Op. pars: C=0.859; sigma=0.0454
## auc.test.mean=0.8439726
##
## [[3]]
## Tune result:
## Op. pars: C=173; sigma=0.00819
## auc.test.mean=0.8462711
```

# PARALLELIZATION

- We use our own package: `parallelMap`
- Setup:

```
> parallelStart("multicore")  
> benchmark(...)  
> parallelStop()
```

- Backends: `local`, `multicore`, `socket`, `mpi` and `batchtools`
- The latter means support for: makeshift SSH-clusters, Docker swarm and HPC schedulers like SLURM, Torque/PBS, SGE or LSF
- Levels allow fine grained control over the parallelization
  - ▶ `mlr.resample`: Job = “train / test step”
  - ▶ `mlr.tuneParams`: Job = “resample with these parameter settings”
  - ▶ `mlr.selectFeatures`: Job = “resample with this feature subset”
  - ▶ `mlr.benchmark`: Job = “evaluate this learner on this data set”

# INTERPRETABLE MACHINE LEARNING

- `iml` - Interpretable Machine Learning -  
<https://github.com/christophM/iml>
- Background
  - ▶ Machine learning has a huge potential
  - ▶ Lack of explanation hurts trusts and creates barrier for machine learning adoption
  - ▶ Interpretation of the behaviour and explanation of predictions of machine learning model with **Interpretable Machine Learning**

# SUPPORTED METHODS

- Model-agnostic interpretability methods for **any** kind of machine learning model
- Supported are
  - ▶ Feature importance
  - ▶ Partial dependence plots
  - ▶ Individual conditional expectation plots
  - ▶ Tree surrogate
  - ▶ Local interpretable model-agnostic explanations
  - ▶ Shapley value

# ONE IML MODEL FOR ALL METHODS I

- Use `iml` package

```
> library(iml)
```

- We use our trained model `mod`
- We need training data from the index vector `train`

```
> mod
```

```
## Model for learner.id=classif.randomForest; learner.class=clas  
## Trained on: task.id = titanic; obs = 872; features = 10  
## Hyperparameters:
```

# ONE IML MODEL FOR ALL METHODS II

- Extract features
- Create IML model


```
> X = dropNamed(train.data, "Survived")  
> iml.mod = Predictor$new(mod, data = X,  
+   y = train.data$Survived, class = 2)
```



# FEATURE IMPORTANCE

- What were the most important features?


```
> imp = FeatureImp$new(iml.mod, loss = "ce")  
> plot(imp)
```



# PARTIAL DEPENDENCE PLOTS

- How does the “passenger class” influence the prediction on average?


```
> pdp = PartialDependence$new(iml.mod, feature = "Pclass")  
> plot(pdp)
```




# LOCAL LINEAR MODELS (LIME)

- Explain a single prediction with LIME

```
> X[1,]  
  
##      Pclass      Sex Age Sibsp Parch      Fare Embarked  farePp de  
## 1         1 female  29     0     0 211.3375          S 211.3375  
  
> lime = LocalModel$new(iml.mod, x.interest = X[1,])  
> plot(lime)
```



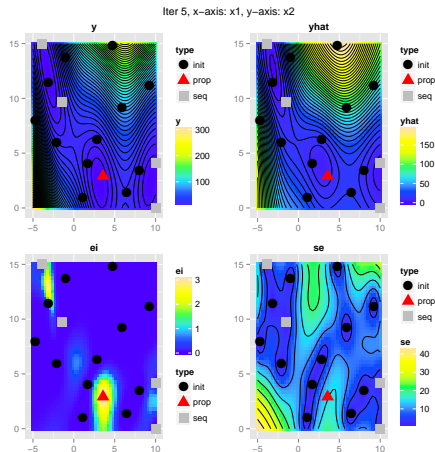
# EXPENSIVE BLACK-BOX OPTIMIZATION



- mlrMB0 - Bayesian Optimization and Model-Based Optimization - <https://github.com/mlr-org/mlrMB0>
- Goal: optimize *expensive black box functions* by *model-based optimization* (aka Bayesian optimization)

# MLRMBO: MODEL-BASED OPTIMIZATION TOOLBOX

- Any regression from mlr
- Arbitrary infill
- Single - or multi-crit
- Multi-point proposal
- Via parallelMap and batchtools runs on many parallel backends and clusters
- Algorithm configuration
- Active research



# MLR-MBO I

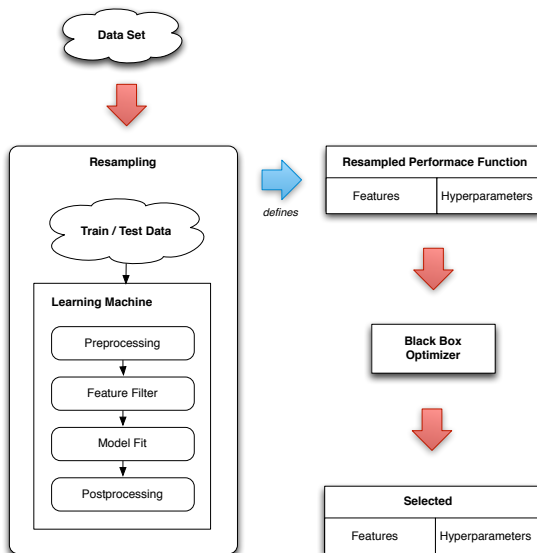
Create an unified interface with the general mlrMBO workflow

1. Define **objective function** and its parameters using the package `smoof`
2. Generate **initial design** (optional)
3. Define mlr' learner for **surrogate model** (optional)
4. Set up a **MBO control** object
5. Start the optimization with `mbo()`

Supported are

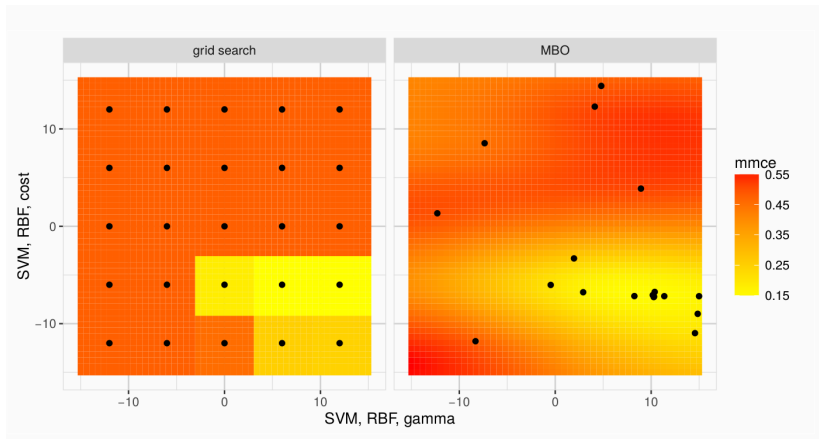
- Efficient global optimization (EGO) of problems with numerical domain and Kriging as surrogate
- Using arbitrary regression models from mlr as surrogates
- Built-in parallelization using multi-point proposals
- Mixed-space optimization with categorical and subordinate parameters, for parameter configuration and tuning
- Multi-criteria optimization

# FROM NORMAL MBO TO HYPERPARAMETER TUNING





# HYPERPARAMETER TUNING



# EXAMPLE: PACKAGES AND CONFIGURATION

```
> library(mlrMBO) # Bayesian Optimization in R  
> library(ParamHelpers) # Objects for parameter spaces  
> library(smoof) # Interface for objective functions  
> set.seed(2)
```

- We run all optimization with VERY few evals to reduce time and log output

```
> iters = 5
```

# EXAMPLE: MIXED SPACE OPTIMIZATION

- Extend our parameter set from Titanic example

```
> par.set = makeParamSet(  
+   makeNumericParam("C", lower = -8, upper = 8,  
+   trafo = function(x) 2^x),  
+   makeNumericParam("sigma", lower = -8, upper = 8,  
+   trafo = function(x) 2^x)  
+ )
```

# OBJECTIVE FUNCTION

- We use our Titanic task
- Create a single objective function

```
> svm = makeSingleObjectiveFunction(name = "svm.tuning",  
+   fn = function(x) {  
+     # remove inactive parameters coded with `NA`  
+     x = x[!vapply(x, is.na)]  
+     lrn = makeLearner("classif.ksvm", par.vals = x)  
+     crossval(lrn, task, iters = 2, show.info = FALSE)$aggr  
+   },  
+   par.set = par.set,  
+   noisy = TRUE,  
+   has.simple.signature = FALSE,  
+   minimize = TRUE  
+ )  
> ctrl = makeMBOControl()  
> ctrl = setMBOControlTermination(ctrl, iters = iters)
```

# MBO LEARNER I

- Parameter set is not a suitable surrogate
- Use random Forest with imputation for non-active parameters

```
> makeMBOlearner(ctrl, svm)

## Learner regr.km from package DiceKriging
## Type: regr
## Name: Kriging; Short name: km
## Class: regr.km
## Properties: numerics,se
## Predict-Type: se
## Hyperparameters: jitter=TRUE,covtype=matern3_2,optim.method=g
```


# MBO LEARNER II

```
> res = mbo(svm, control = ctrl)
> print(res)

## Recommended parameters:
## C=-0.305; sigma=-2.19
## Objective: y = 0.202
##
## Optimization path
## 8 + 5 entries in total, displaying last 10 (or less):
##           C          sigma          y dob eol error.message exec.
## 4  -7.7416820 -6.333102 0.3819584    0  NA          <NA>
## 5   4.9360370  3.099967 0.3307725    0  NA          <NA>
## 6  -0.8946519 -3.522210 0.2054917    0  NA          <NA>
## 7  -4.4789734  4.361640 0.3819677    0  NA          <NA>
## 8   6.8105644 -4.292903 0.2177382    0  NA          <NA>
## 9   4.2962856 -2.585426 0.2207811    1  NA          <NA>
## 10  7.9947749 -3.508953 0.2360599    2  NA          <NA>
## 11  1.8919236 -4.051119 0.2207683    3  NA          <NA>
## 12 -0.3054272 -2.186221 0.2024371    4  NA          <NA>
## 13 -0.1544471 -2.866955 0.2024441    5  NA          <NA>
## error model train time  prop type propose time
```

# RESULTS

```
> op = as.data.frame(res$opt.path)
> plot(cummin(op$y), type = "l", ylab = "mmce",
+      xlab = "iteration")
```



# REFERENCES

- mlrMBO Paper on arXiv (under review)  
<https://arxiv.org/abs/1703.03373>
- Bischl, Wessing et al: *MOI-MBO: Multiobjective infill for parallel model-based optimization*, LION 2014
- Horn, Wagner, Bischl et al: *Model-based multi-objective optimization: Taxonomy, multi-point proposal, toolbox and benchmark*, EMO 2014



# MLRCPO I

- mlrCPO - Composable Preprocessing Operators for mlr - <https://github.com/mlr-org/mlrCPO>

```
> library(mlrCPO)
```

- Preprocessing operations (e.g. imputation or PCA) as R objects with their own hyperparameters

```
> operation = cpoScale()  
> print(operation)  
## scale(center = TRUE, scale = TRUE)
```

# MLRCPO II

- Objects are handled using the “piping” operator %>%:
- Composition:

```
> imputing.pca = cpoImputeMedian() %>% cpoPca()
```

- Application to data

```
> task %>% imputing.pca
```

- Combination with a Learner to form a machine learning pipeline

```
> pca.rf = imputing.pca %>%  
+   makeLearner("classif.randomForest")
```

# MLRCPO EXAMPLE: TITANIC I

The feature engineering and preprocessing steps done on the Titanic dataset, using mlrCPO:

```
> # Add interesting columns
> newcol.cpo = cpoAddCols(
+   farePp = Fare / (Parch + Sibsp + 1),
+   deck = stri_sub(Cabin, 1, 1),
+   side = {
+     digit = stri_sub(Cabin, 3, 3)
+     digit = suppressWarnings(as.numeric(digit))
+     c("port", "starboard")[digit %% 2 + 1]
+   })
```

## MLRCPO EXAMPLE: TITANIC II

```
> # drop uninteresting columns
> dropcol.cpo = cpoSelect(names = c("Cabin",
+   "Ticket", "Name"), invert = TRUE)
>
> # impute
> impute.cpo = cpoImputeMedian(affect.type = "numeric") %>%
+   cpoImputeConstant("__miss__", affect.type = "factor")
```

# MLRCPO EXAMPLE: TITANIC III

```
> train.task = makeClassifTask("Titanic", train.data,  
+   target = "Survived")  
>  
> pp.task = train.task %>>% newcol.cpo %>>%  
+   dropcol.cpo %>>% impute.cpo
```

- Advantage: Different preprocessing steps can be tried by preparing different CPO objects (→ “strategy pattern”).

# TRANSFORMATION OF NEW DATA

- New data (e.g. for testing, prediction) must also be preprocessed, in same order and with same hyperparameters
- Preprocessing parameters (e.g. PCA matrix) should only depend on training data
- Use `retrafo()` to get retrafo information to use on test data
- Object of type `CPOTrained`, behaves very similar to `CP0`

```
> # get retransformation
> ret = retrafo(pp.task)
> # can be applied to data using the %>>% operator,
> # just as a normal CP0
> pp.test = test.data %>>% ret
```

# COMBINATION WITH LEARNERS

- Attach one or more CPO to a Learner to build machine learning pipelines
- Automatically handles preprocessing of test data

```
> learner = newcol.cpo %>>% dropcol.cpo %>>%  
+   impute.cpo %>>% makeLearner("classif.randomForest",  
+   predict.type = "prob")  
>  
> # the new object is a "CPOLearner", subclass of "Learner"  
> inherits(learner, "CPOLearner")  
  
## [1] TRUE  
  
> # train using the task that was not preprocessed  
> ppsmod = train(learner, train.task)
```

# TUNING WITH MLRCPO I

- CPO hyperparameters can be tuned in combination with Learner parameters
- Tuning can be done using `tuneParams()` function from `mlr`

```
> class(learner)[[1]]  
  
## [1] "CPOLearner"  
  
> ps = makeParamSet(  
+   makeIntegerParam("ntree", lower = 1, upper = 500),  
+   makeIntegerParam("mtry", lower = 1, upper = 10)  
+ )
```



# TUNING WITH MLRCPO II

```
> tuneParams(learner, train.task, cv3, par.set = ps,  
+   control = makeTuneControlRandom(maxit = 10L),  
+   measures = mlr::auc)  
  
## Tune result:  
## Op. pars: ntree=460; mtry=3  
## auc.test.mean=0.8604592
```

# MLRCPO III

- `listCPO()` to show available CPOs
- Currently 69 CPOs, and growing: imputation, feature type conversion, target value transformation, over/undersampling, ...
- “cbind” CPO combines different preprocessing outputs of the same data

```
> scale = cpoSelect(pattern = "Fare", id = "first") %>%  
+   cpoScale(id = "scale")  
> scale.pca = scale %>% cpoPca()  
> cbinder = cpoCbind(scale, scale.pca, cpoSelect(  
+   pattern = "Age", id = "second"))  
> result = train.data %>% cbinder  
> result[1:3, ]
```

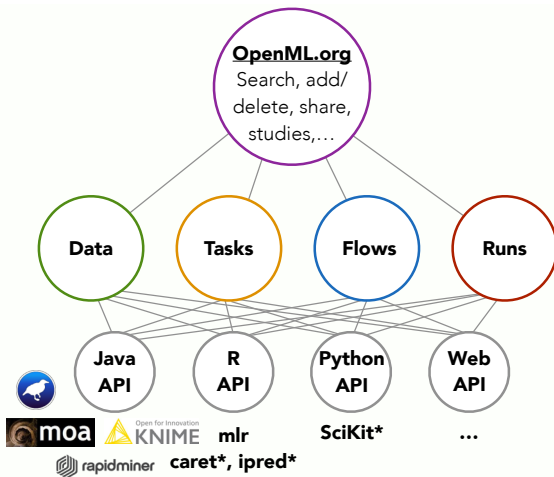
```
##      Fare      PC1      Age  
## 1 3.768222 3.768222 29.0000  
## 2 2.512035 2.512035  0.9167  
## 4 2.512035 2.512035 30.0000
```

- CPO “multiplexer” enables tuning over different distinct preprocessing operations
- Custom CPOs can be created using `makeCPO()`
- Further documentation in the vignettes:

```
> vignette("a_1_getting_started")
```

# OPENML

Main idea: Make ML experiments reproducible, computer-readable and allow collaboration with others.



# OPENML R-PACKAGE

<https://github.com/openml/r>

## TUTORIAL

- Caution: Work in progress

## CURRENT API IN R

- Explore and Download data and tasks
- Register learners and upload runs
- Explore your own and other people's results

# OPENML ACCOUNT

- Install the openML package and either farff or RWeka

```
> library("OpenML")
```

- You need an openML API key to talk to the server
- Create an account on <https://www.openml.org/register>

```
> setOMLConfig(apikey = "c1994bdb7ecb3c6f3c8f3b35f4b47f1f")  
>  
> # Permanently save your API disk to your config file  
> saveOMLConfig(apikey = "c1994...47f1f", overwrite=TRUE)
```

- Find your own API key in account settings API Authentication

# OPENML DATA AND TASKS I

- You can access all datasets or tasks

```
> datasets = listOMLDataSets()  
> datasets[1:3, c(1,2,11)]
```

```
##   data.id      name number.of.features  
## 1        2   anneal                 39  
## 2        3 kr-vs-kp                 37  
## 3        4    labor                 17
```

```
> tasks = listOMLTasks()  
> tasks[1:3, 1:4]
```

```
##   task.id          task.type data.id      name  
## 1        2 Supervised Classification      2   anneal  
## 2        3 Supervised Classification      3 kr-vs-kp  
## 3        4 Supervised Classification      4    labor
```

# OPENML DATA AND TASKS II

- Search for data on <https://www.openml.org/home>

The screenshot shows the OpenML website interface. At the top, there is a green header with the OpenML logo and the search term 'titanic'. Below the header, on the left, is a sidebar with navigation options: Data (8), Task, Flow, Run, Study, Task type, Measure, and People. The main content area displays 8 search results for 'titanic'. The results are filtered to show only active datasets (public or shared with you). The results are sorted by best match. The first three results are visible:

Dataset Name	Description	Stats
Titanic (1)	The original Titanic dataset, describing the survival status of individual passengers on the ...	★ 0 runs ♥ 0 likes 📄 2 downloads 📶 2 reach ⚡ 1 impact 1309 instances - 14 features - 2 classes - 3855 missing values
Titanic (2)	PMLB version of the Titanic dataset, which only uses 3 features. See version 1 for the com...	★ 3 runs ♥ 0 likes 📄 0 downloads 📶 0 reach ⚡ 4 impact 2201 instances - 4 features - 2 classes - 0 missing values
analcatdata_broadwaymult (...)	Binarized version of the original data set (see version 1). The multi-class target feature is ...	★ 694 runs ♥ 0 likes 📄 5 downloads 📶 5 reach ⚡ 5 impact 285 instances - 8 features - 2 classes - 27 missing values



# OPENML TITANIC DATASET

- We download the Titanic dataset from OpenML

```
> listOMLDataSets(data.name = "titanic")[, 1:5]

##   data.id   name version status format
## 1   40704 Titanic      2 active  ARFF
## 2   40945 Titanic      1 active  ARFF

> titanic = getOMLDataSet(data.id = 40945L)
```

# OPENML TITANIC TASK

- We also can directly load the Titanic classification task

```
> listOMLTasks(data.name = "titanic")[1:2, 1:4]

##      task.id                task.type data.id    name
## 1   145769                Clustering  40704 Titanic
## 2   146230 Supervised Classification  40704 Titanic

> titanic.task = getOMLTask(task.id = 146230)
> titanic.task

##
## OpenML Task 146230 :: (Data ID = 40704)
##   Task Type           : Supervised Classification
##   Data Set            : Titanic :: (Version = 2, OpenML ID =
##   Target Feature(s)   : class
##   Estimation Procedure : Stratified crossvalidation (1 x 10 f
##   Evaluation Measure(s): precision
```

# OPENML AND MLR

- We can use OpenML and mlr together
- Use mlr for learner and use the task that we've got from OpenML

```
> lrn = makeLearner("classif.randomForest", mtry = 2)
> run.mlr = runTaskMlr(titanic.task, lrn)
> run.mlr$bmr$results

## $Titanic
## $Titanic$classif.randomForest
## Resample Result
## Task: Titanic
## Learner: classif.randomForest
## Aggr perf: ppv.test.join=0.7692308,timetrain.test.sum=3.94000
## Runtime: 4.17194
```

# OPENML UPLOAD

- You can upload your own data sets to OpenML
- Three steps are necessary
  1. `makeOMLDataSetDescription`: create the description object of an OpenML data set
  2. `makeOMLDataSet`: convert the data set into an OpenML data set
  3. `uploadOMLDataSet`: upload the data set to the server
- We can upload our Titanic data set to OpenML

```
> titanic.desc = makeOMLDataSetDescription(name = "titanic",  
+   description = "Titanic data set ...")  
>  
> titanic.data = makeOMLDataSet(desc = titanic.desc,  
+   data = data, target.features = "Survived")  
>  
> # titanic.id = uploadOMLDataSet(titanic.data)
```

# THERE IS MORE ...

- Regression, Clustering and Survival analysis
- Cost-sensitive learning
- Multi-Label learning
- Imbalancy correction
- Wrappers
- Bayesian optimization
- Multi-criteria optimization
- Ensembles, generic bagging and stacking
- ...

# WE ARE WORKING ON

- Even better tuning system
- More interactive and 3D plots
- Large-Scale learning on databases
- Time-Series tasks
- Large-Scale usage of OpenML
- auto-mlr
- ...

# MLR CONTRIBUTION

- Write an issue on Git
- We are founding an association - **Machine Learning in R e.V**  
subscribe for updates [contact.mlr.org@gmail.com](mailto:contact.mlr.org@gmail.com)

Thanks!